

# Analyzing SQL Statements for Performance

---



**David Berry**

@davidcberry13 [buildingbettersoftware.blogspot.com](http://buildingbettersoftware.blogspot.com)



# Module Outline



**Generating an execution plan**

**Reading an execution plan**

**Viewing execution statistics**

**Rewriting SQL statements for performance**

**Common execution plan operations**



# Statement Performance Comparison

Metric	Without Index	With Index
Data Access Operation	Clustered index scan	Index seek
Statement Cost	14.921	0.189
Logical Reads	12,116	305
CPU Time (ms)	141 ms	< 1 ms
Elapsed Time (ms)	136 ms	< 1 ms





Rewriting a statement sometimes improves performance

Subqueries can be rewritten as joins

Joins can be rewritten as subqueries

EXISTS and NOT EXISTS are best in some scenarios



# SQL Server Data Access Operations

Operation	Description
Clustered Index Scan	Reads all of the rows stored in a table stored as a clustered index
Table Scan	Reads all the rows in a table that is stored as a heap structure
Clustered Index Seek	Traverses the tree structure of a table stored as a clustered index to find the needed row(s)
Index Scan	Reads all of the key values of an index to find the matching data
Index Seek	Traverses the tree structure of an index to find the matching index keys



# SQL Server Join Operations

## Operation

## Description

---

Nested Loops Join

For each value in the first data set, SQL Server loops through the second data set looking for matches

Merge Join

Used to join two data sets that are already sorted using the same key. A row from each source is obtained. If the rows match they are joined. If the rows do not match, the lower value row is discarded and a new row is obtained from that source.

Hash Match

A hashtable of the smaller data set is created, then SQL Server loops through the larger data set probing the hashtable for matching values. Used when two large data sets must be joined



# Statement Tuning Options

Add an index

Rewrite subqueries/joins/EXISTS clauses

Evaluate if statement can be more selective



Getting an **execution plan** helps you identify the **most costly operations** within your statement so you can **tune those operations**





Up Next

---

Building Effective Database Indexes

