

Building Effective Indexes



David Berry

@davidcberry13 buildingbettersoftware.blogspot.com



Module Outline



Index terminology refresher

What columns should I index?

Why column order matters in an index?

Index selectivity

Include columns and covering indexes

Functions in the WHERE clause and indexes

Over-indexing

SQL Server's index recommendations

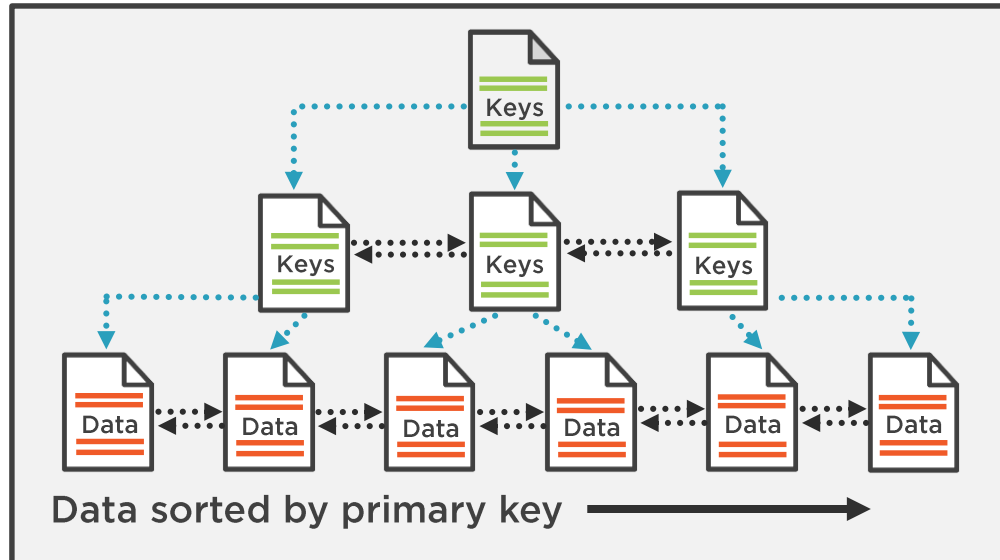


Index Terminology Refresher



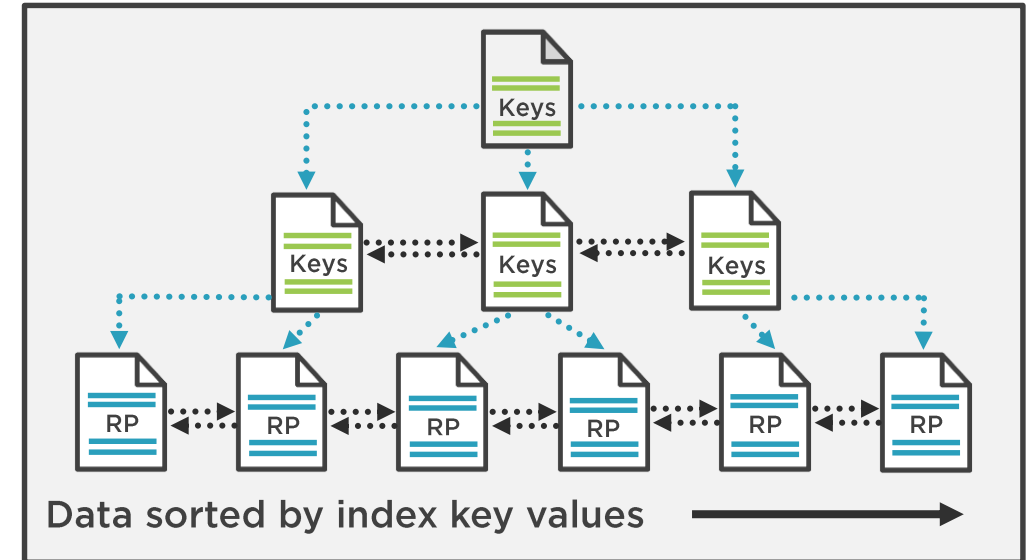
Clustered vs. Non-clustered Indexes

Clustered Index



- Stores the data for the table
- Typically organized by the primary key
- Only one allowed per table


Non-clustered Index



- Structure used for all other indexes
- Data organized by index key
- Contains pointers to matching rows
- Multiple allowed per table


Scan vs. Seek Operations

Index Scan Operation



SQL Server reads all data in the index

Index Seek Operation



SQL Server traverses the index using the tree structure




Indexing Strategy

By correctly indexing columns up front, you can save yourself from most performance problems later




Add Indexes for WHERE Clause Criteria

```
SELECT *  
FROM Students  
WHERE LastName = 'Howard'  
AND FirstName = 'Emily'
```



Index on LastName, FirstName

```
SELECT *  
FROM Students  
WHERE Email =  
    'Ehoward@someuniversity.edu'
```




Index on Email

Primary key columns are indexed
by default in SQL Server



Index Foreign Key Columns



Helps speed up join operations

Applications often query tables by foreign key values

Applying Index Column Order Lessons

How Data is Queried

```
SELECT * FROM Students  
WHERE LastName = ?????
```

```
SELECT * FROM Students  
WHERE LastName = ????? AND FirstName = ?????
```

```
SELECT * FROM Students  
WHERE LastName = ????? AND FirstName = ?????  
AND State = ?????
```

```
SELECT * FROM Students  
WHERE State = ????? AND City = ?????
```

Index IX_Students_NameSearch

LastName	FirstName	State
----------	-----------	-------

Index IX_Students_LocationSearch

State	City
-------	------



Primary Key PK_Students (Clustered Index)

StudentId

Created by default. Supports lookups, updates by student id

Index IX_Students_NameSearch

LastName

FirstName

State

Lookup students by their name

Index IX_Students_LocationSearch

State

City

Lookup students by their home state and city

Index IX_Students_Email

Email

Lookup students by their email address



Index Selectivity

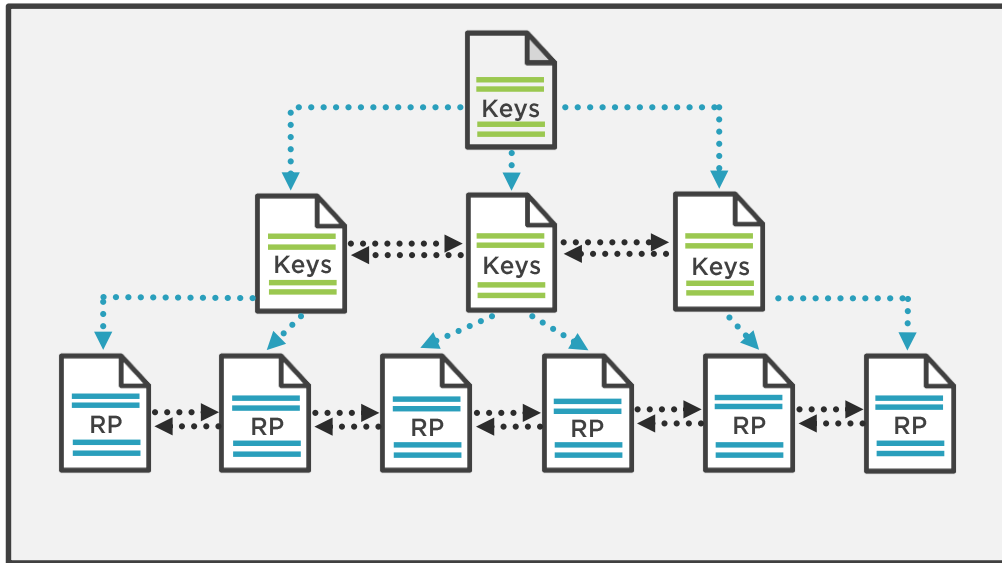
A measure of how many or how few rows correspond to each index key value



For SQL Server to use an index, it must be selective

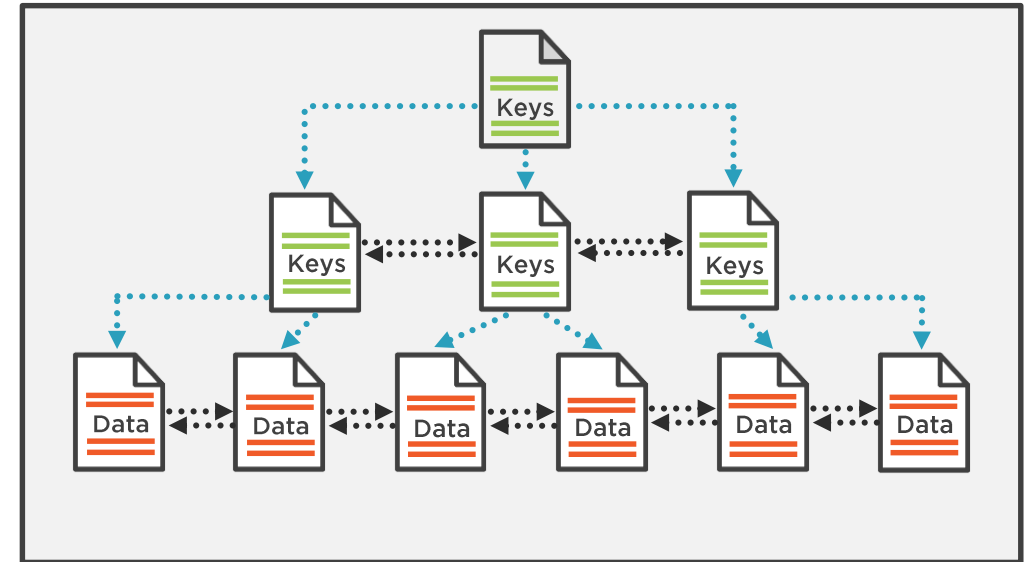
Index Operations

Index



SQL Server looks up
matching index keys, gets
row identifiers

Table



SQL Server looks up
corresponding rows by their
row identifier

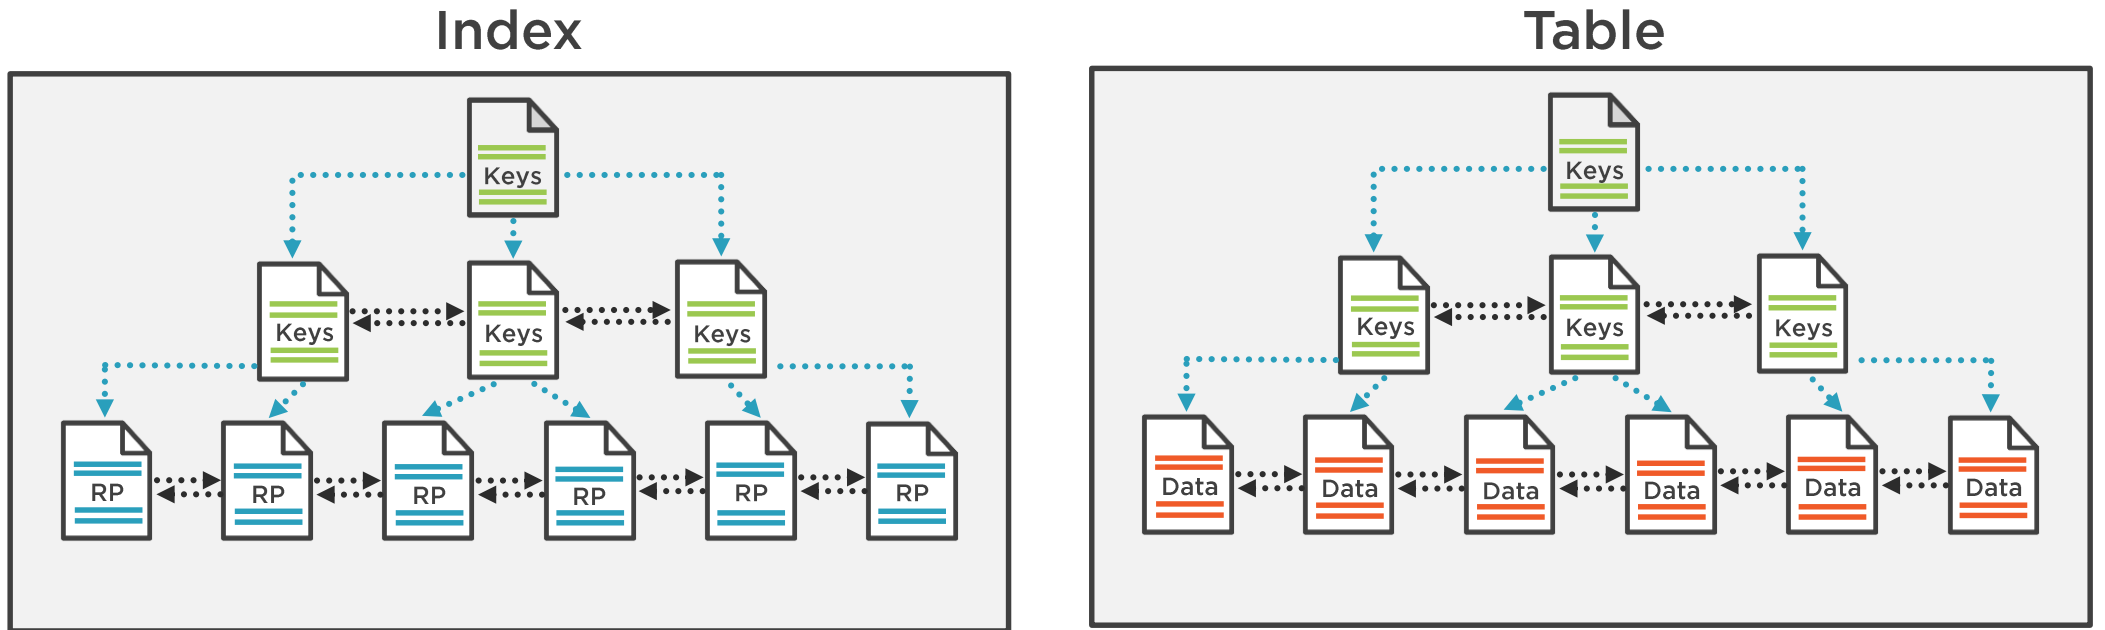


Effects of Index Selectivity

	Index Operation	Table Lookup Operation	Overall Cost
Index is Selective	Few rows for each matching key	Few rows to lookup	Low. Very few pages to read
Index is not Selective	Many rows for each matching key	Many rows to look up	High. Many pages to read



Index Lookup Operations



SQL Server finds matching keys

Index is
Selective

Index is not
Selective



Selective indexes help SQL Server locate the exact data needed with minimal effort



Statement Performance Comparison

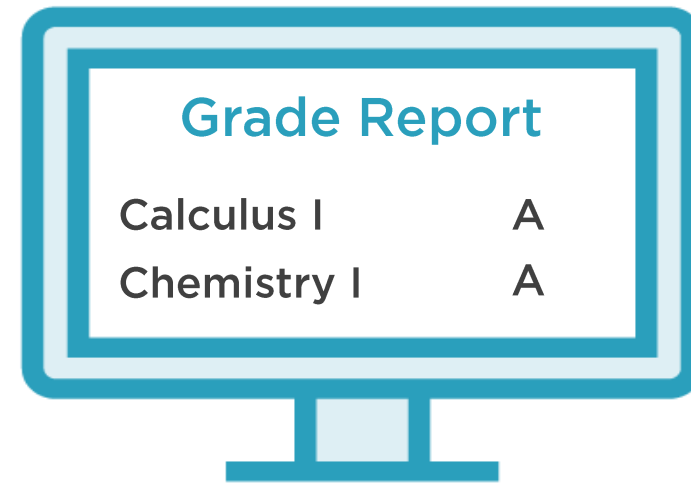
Metric	Chosen by SQL Optimizer	Index Usage Forced with Hint
Data Access Method	Table scan	Index seek
Statement Cost	2.903	5.826
Logical Reads	3747	7189



Queries Often Follow Foreign Keys



Student ID = 12345



```
SELECT c.DepartmentCode,  
       c.CourseNumber, c.CourseTitle,  
       ce.Grade,  
FROM CourseEnrollments ce  
    ...  
WHERE ce.StudentId = 12345
```



Covering Index

SQL Server is able to find all the data for the query in the index itself

This avoids needing a key lookup operation to find the data row in the table



Consider a **covering index** when you only need to add **one or two columns** as include columns to give a **performance boost** to a key query



Question:

If indexes are so great, then why not create an index on every column?



Answer:

Indexes require maintenance, so you only want to create indexes that are actually used by your SQL statements




Table Data

Id	FirstName	LastName	Email
101	Mark	Gilbert	mgilbert@...
102	Paul	Wilson	pwilson@...
103	Erica	Griffin	egriffin@...
105	Peter	Wright	pwright@...

DML Statement
(INSERT, UPDATE or DELETE)


Index update

Index




Index update

Index



Index update

Index



Regard Indexes as an Investment

Indexes that are frequently used repay this investment many times

Indexes that are seldom or never used add cost without adding value



SQL Server **Dynamic Management Views** provide access to a rich set of performance data, including **index usage data**



Finding Unused Indexes Using DMVs

```
SELECT
    OBJECT_NAME(s.object_id) AS TableName,
    i.name As IndexName,
    i.type_desc as IndexType,
    user_seeks + user_scans + user_lookups As TotalUsage,
    user_seeks,
    user_scans,
    user_lookups,
    user_updates
FROM sys.dm_db_index_usage_stats s
RIGHT OUTER JOIN sys.indexes i
    ON s.[object_id] = i.[object_id]
    AND s.index_id = i.index_id
WHERE s.database_id = DB_ID()
    AND i.name IS NOT NULL
    AND OBJECTPROPERTY(s.[object_id], 'IsMsShipped') = 0
ORDER BY s.object_id, s.index_id
```



SQL Server Index Recommendations

Analyze the recommendation first
before creating

SQL Server can be overly aggressive



SQL Server's recommendations are based on looking at one individual SQL statement

You want to view these recommendations in the larger context of all of your statements for a given table



Module Review



Index Columns Used in WHERE Clauses

```
SELECT *  
  FROM Students  
 WHERE LastName = ????  
    AND FirstName = ????;
```

```
SELECT *  
  FROM Students  
 WHERE Email = ????;
```

```
SELECT *  
  FROM Students  
 WHERE State = ????  
    AND City = ????;
```

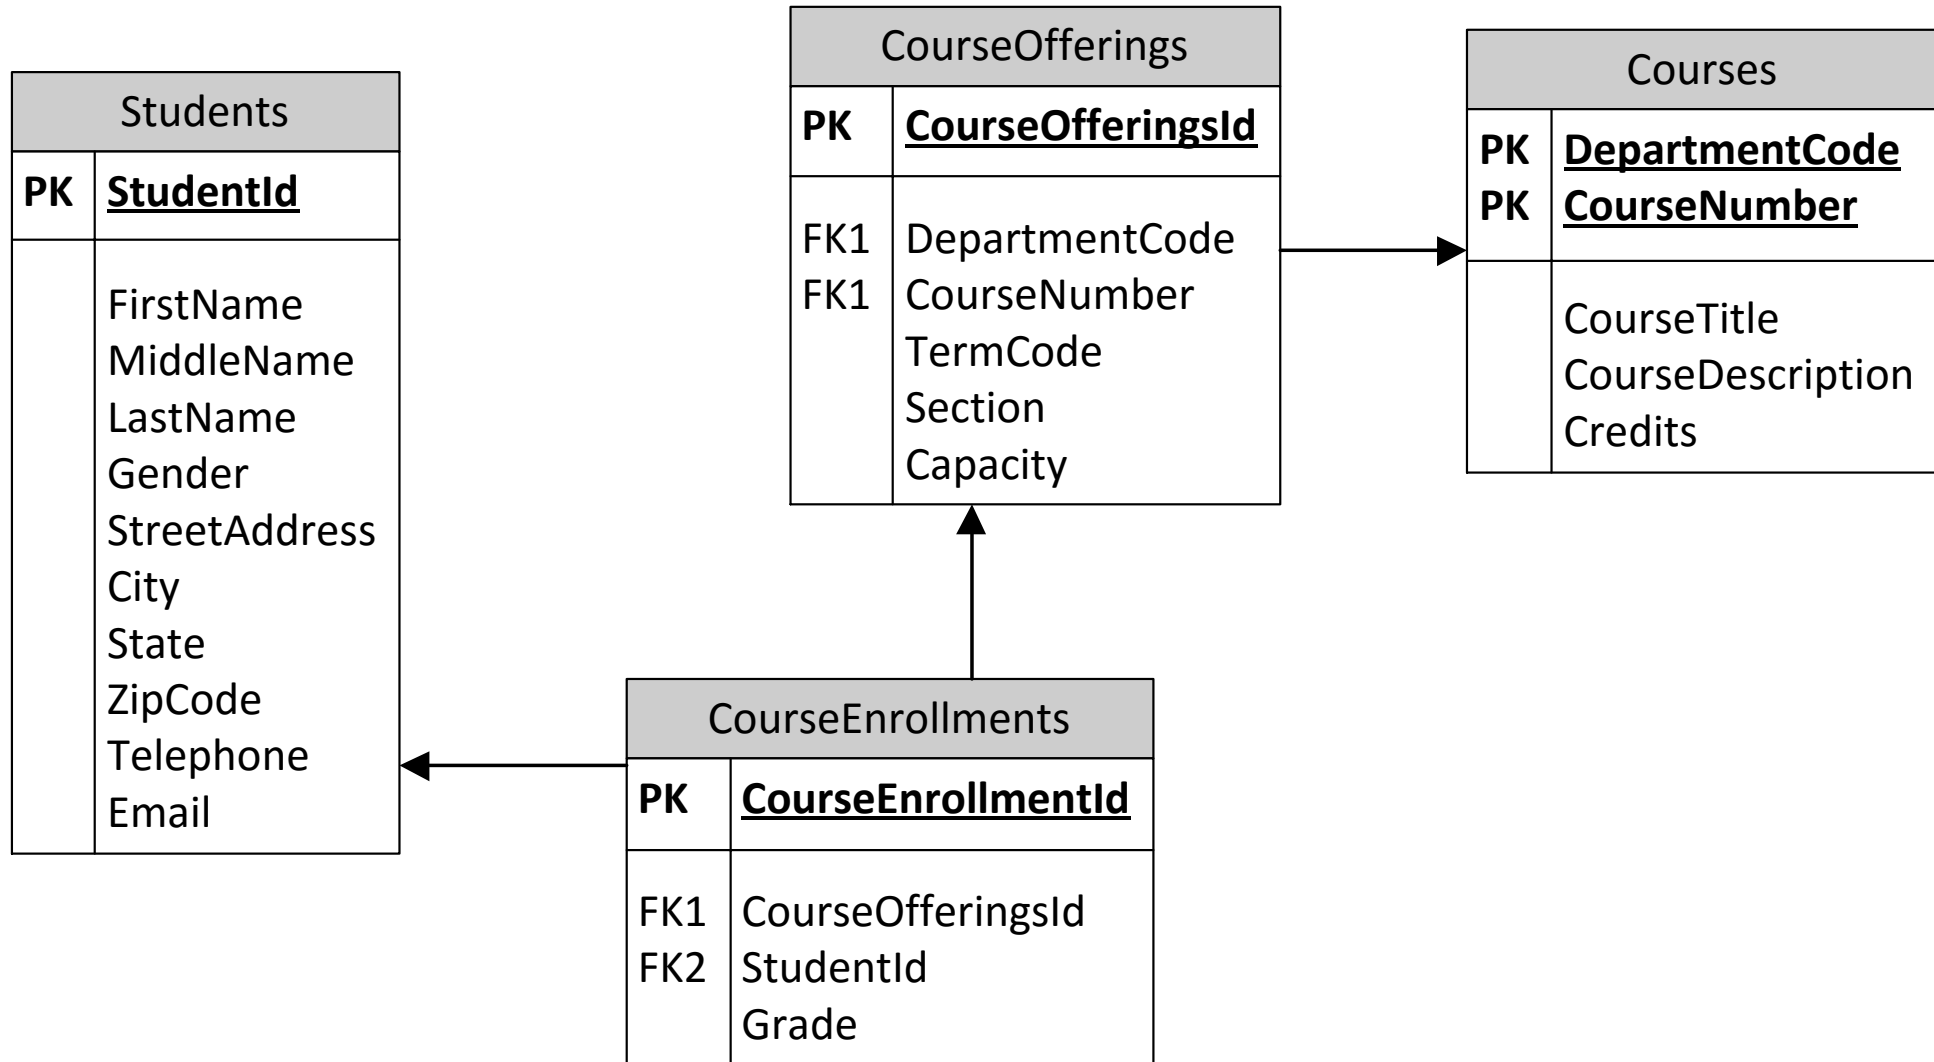
```
SELECT *  
  FROM Applicants  
 WHERE Email = ????;
```

```
SELECT *  
  FROM Applicants  
 WHERE LastName = ????  
    AND FirstName = ????;
```

```
SELECT *  
  FROM CourseOfferings  
 WHERE TermCode = ????  
    AND DepartmentCode = ????  
    AND CourseNumber = ????;
```



Index Foreign Key Columns



Column Order in Indexes Matters

```
SELECT *  
  FROM Students  
 WHERE LastName = '????'  
        AND FirstName = '????'
```

```
SELECT *  
  FROM Students  
 WHERE LastName = '????'
```

```
SELECT *  
  FROM Students  
 WHERE LastName = '????'  
        AND DegreeId = '????'
```

Index



LastName	FirstName	DegreeId
----------	-----------	----------



Index Selectivity

Index

State

➔
$$\frac{120,000 \text{ rows}}{50 \text{ unique values}} = 2400 \text{ rows/index key}$$

**Poor
Selectivity**

Index

State	City
-------	------

➔
$$\frac{120,000 \text{ rows}}{13,000 \text{ unique values}} = 9.2 \text{ rows/index key}$$

**Good
Selectivity**

Index

LastName	FirstName
----------	-----------

➔
$$\frac{120,000 \text{ rows}}{107,000 \text{ unique values}} = 1.1 \text{ rows/index key}$$

**Good
Selectivity**



-- This statement will not use an index

```
SELECT *  
    FROM Students  
    WHERE LastName = '%Harris%' AND FirstName = '%Christy%'
```

-- This statement will use an index

```
SELECT *  
    FROM Students  
    WHERE LastName = 'Harris%' AND FirstName = 'Christy%'
```

-- But beware of statements like this where you don't

-- provide much information. It probably will not use an index

```
SELECT *  
    FROM Students  
    WHERE LastName = 'Ha%' AND FirstName = 'Ch%'
```

Functions in the WHERE Clause

-- Statements like this will not use a normal index

```
SELECT *  
  FROM Students  
 WHERE SOUNDEX(LastName) = SOUNDEX('SMITH')  
    AND SOUNDEX(FirstName) = SOUNDEX('CHRIS');
```

-- But we can create a computed column and an index over the computed column if we need this functionality

```
ALTER TABLE Students ADD LastNameSoundex AS Soundex(LastName);  
ALTER TABLE Students ADD FirstNameSoundex AS Soundex(FirstName);
```

```
CREATE INDEX IX_Students_PhoneticName  
  ON Students (LastNameSoundex, FirstNameSoundex);
```



Up Next

Finding Performance Bottlenecks in SQL Server

